

DOCS pitch deck

Read access to desk/reports/ wasn't granted, so I'm delivering the deck inline rather than writing it to a file. Say the word if you want it saved to desk/reports/DOCS_PITCH_DECK_20260821.md.

PITCH — DOCS — VERDICT: PROPOSED — PENDING ADJUDICATION

No entry today. Both benches land on no-position; the disagreement is about why, and both of them got the stock price wrong.

What we believe

Doximity is a profitable, debt-free advertising and hiring platform for doctors. It has \$687.8M of cash and securities, no borrowings, and generated \$39.6M of free cash in the June quarter. It is not in any danger.

The problem is what you pay for it and what you get. Revenue growth has slowed to roughly +7% year over year (\$156.6M in the June quarter against \$145.9M a year earlier), and the company guided the September quarter to \$170–171M, which is only about +1% against last year. At the same time, the company now pays its employees about \$147M a year in stock — roughly 22% of all revenue — and that cost is real even though it doesn't show up in the profit number the company highlights.

If you treat stock pay as the expense it is, you are paying about **25 times** the company's true annual operating profit for a business growing 7%. That is expensive.

The thing that makes expensive dangerous here: Doximity has two classes of stock, and the founders' shares carry extra votes. That means no outside investor can force change, and no acquirer can buy the company without the founders' consent. For most richly-priced, asset-light businesses, "someone could just buy it" is the floor under the price. Here there is no such floor.

So: a good business, priced for a better one, with the usual safety net cut away. We pass, and we say exactly what would change our minds.

What the market believes

The market does not have one view — it has two incompatible ones, which is unusual and is the only reason this is still an open file.

- The stock is **\$25.89** (broker feed, 2026-08-21). It is **down 66% from its 52-week high of \$76.50** and **up 51% from its 52-week low of \$17.16**. Something already broke; the argument is over whether it's fixed.

- On revenue, everyone agrees. Wall Street's average forecast for this fiscal year is about **\$670.5M**; Truist just raised its number to **\$673M**; the company itself guided **\$676M**. Our own work landed at \$676M. There is no revenue disagreement to trade.
- On what it's worth, the disagreement is enormous. **Wells Fargo downgraded the stock and put a \$18 target on it. Needham raised its target to \$41.** That is a 2.3-to-1 spread at a \$25.89 stock. Translated into the measure we care about, the street band is **16 times to 42 times** post-stock-pay operating profit — the analysts are not arguing about the business, they are arguing about the multiple.

In plain terms: the market believes the revenue story is settled and the valuation story is not.

Why we might have edge

Three narrow places, and we should not overstate them:

1. **We expense stock pay; most published models don't.** The company's own headline profit figure (\$319M for the year) ignores \$147M of stock compensation. Nobody has to agree with our adjustment, but it is the difference between "13 times profit" (cheap-ish) and "25 times profit" (expensive). Since the street's fight is entirely about the multiple, the adjustment sits exactly on the contested variable.
2. **The vote structure appears in nobody's target price.** Neither our original brief nor, as far as we can tell from the published research, the street's, prices the absence of a takeover floor. For a business at 6.3 times sales, that is the whole downside case.
3. **We have a cleaner forward alarm than a target price.** From the company's own annual guidance we can back out what the December quarter must deliver, and check it against a single line in the November press release. That is a date-stamped, falsifiable test, not an opinion.

Why it might be priced in

Stated honestly, this is the stronger side of the ledger right now:

- **The bear case is already published.** Wells Fargo's \$18 is roughly our own "this is what it's worth if you expense the stock pay" number. We are not early to a view; we are agreeing with a downgrade that already printed at a \$25.89 tape.
- **Two of our three original complaints don't survive checking.** We said growth was structurally breaking; the weak September quarter is lapping the strongest quarter in the company's history, so the two-year comparison is flat, not collapsing. We said stock pay was ramping; it went \$21.9M → \$29.5M → \$33.5M → \$36.7M → \$36.8M per quarter — the ramp ended a quarter ago. What's left is "the level is high," which is a static fact the market has had four quarters to absorb.
- **Margins are guided flat, not falling.** The September guide implies a 47.5% margin, the full year 47.2%, the back half 46.8%. The step down from 54.7% already happened. "Damage arriving" was wrong; "damage arrived and stayed" is right, and arrived damage is usually priced.

- **The share count is shrinking, not growing.** Diluted shares went from 201.2M to 191.2M, down 5.0%. Buybacks are more than absorbing the stock pay. The dilution complaint, as an ownership matter, is backwards.

What we checked ourselves

Including — prominently — where our own benches got it wrong.

Claim	How checked	Found	Confirmed?
Stock price is \$25.21, “the post-earnings pop has faded 3.2%” (asserted by both benches as a correction to the brief)	Read the evidence pack’s own tape line and re-derived it: $\text{high } \$76.50 \times (1 - 0.662) = \25.86 ; $\text{low } \$17.155 \times 1.509 = \25.89	Both benches are wrong. The pack says \$25.89 , 50.9% off the low. \$25.21 appears nowhere in the pack. The brief’s original \$26.035 was within 0.6%; the “fade” is 0.6%, not 3.2%	NO — both benches erred
Enterprise value \$4.13B; $6.1 \times \text{sales} / 13.0 \times \text{profit} / 24.0 \times \text{after stock pay}$ (red bench, presented as <i>superseding</i> the brief)	Recomputed at the real pack price: $\$25.89 \times 191.2\text{M shares} = \4.950B , less \$687.8M net cash = \$4.262B → 6.3× sales, 13.4× profit, 24.8× after stock pay	Red “corrected” the brief to a number that was less accurate. The brief’s original $6.3 / 13.4 / 24.9$ was right. The direction of the case is unchanged; the correction was noise dressed as rigor	NO — red erred; brief vindicated
“Two-year growth is accelerating, so the slowdown is an illusion” (red bench, tagged FATAL)	Blue extended the same calculation across every quarter in the pack instead of the last two	Series runs +34.5%, +48.3%, +36.8%, +23.6%, +24.6% . It halved, then flattened. Red read a rise that exists only between the final two points	NO — red erred; blue’s correction confirmed
“Stock pay is up 68%, margin damage arriving” (original brief)	Pulled the raw SEC data series and differenced it: SEC XBRL , ShareBasedCompensation , CIK 1516513	$\$21.9 \rightarrow \$29.5 \rightarrow \$33.5 \rightarrow \$36.7 \rightarrow$ \$36.8M : up 0.2% in the latest quarter. The increase happened inside last fiscal year. Also kills the brief’s shortcut of multiplying one quarter by four — that understated the actual year by 28%	Brief was wrong; red’s kill confirmed
“A November December-quarter guide below \$193M kills the growth ramp” (blue bench’s headline alarm)	Checked blue’s own arithmetic: it said the December quarter carries ~53% of the second half, and the second half is \$348.9M	53% × \$348.9M = \$184.9M, not \$195M. To reach \$195M you need a 55.9% share. Neither bench published the seasonal split it used, and last year’s December-quarter revenue isn’t in the pack	NO — blue erred; threshold must be recomputed before arming
Margins guided flat, not compressing	Company guidance math: Sept quarter $\$80.5\text{--}\81.5M on $\$170\text{--}\$171\text{M} = 47.5\%$; full year $\$319\text{M}/\$676\text{M} = 47.2\%$; implied back half $\$163.2\text{M}/\$348.9\text{M} = 46.8\%$	Confirmed. Cross-checks against the 8-K figures	YES

Claim	How checked	Found	Confirmed?
Share count is shrinking	Pack data: 201.158M → 191.2M diluted, -5.0%; zero debt	Confirmed. Note the 191.2M is a weighted average from the quarterly report, not the cover-page share count, which neither bench pulled	YES, with a small precision gap
Two share classes remove the takeover floor	Filing structure shows Class A and Class B common with differing votes	Confirmed, and it is the single most important thing in this file. Our own original brief missed it entirely — this is red bench's best work	YES
Customers are paying more slowly (deferred revenue drift)	Deferred revenue \$109.06M vs \$117.42M a year ago = -7.1%, against revenue +7.3%	Real, but not widening: implied billings of \$159.6M are up ~7.1%, matching revenue, and the balance is rebuilding off a March trough. Caveat: red used only the current-portion line and never checked for a long-term portion	YES, but not decisive
The company discloses contracted future revenue	Queried the standard SEC tag	Returns nothing — Doximity does not report this figure, unusual for a subscription business. Not an accusation; it is a hole where a check should be	N/A — not disclosed
Hiring is slowing (would support “they’re just repricing to keep staff”)	Job-posting scrape, San Francisco and Irving TX	Irving sample came back contaminated with a Deloitte posting . Red rejected Irving for that reason, then accepted San Francisco from the same scraper without testing it. One dirty sample indicts the tool	NO — unusable as run; red erred in accepting half of it
Doctor sentiment / app ratings decaying	App-review checker	Lifetime rating 4.824, no decay — but the recent-period comparison, which is the actual test, never returned. Red logged this as a clean pass; it was a partial run	NO — logged clean, actually partial
Insiders selling into the bounce	Four Form 4s (Aug 17) and two Form 144s (Aug 13, Aug 7)	Neither bench opened a single document . Sizes and whether the sales were pre-scheduled are unread. This was the cheapest available evidence in the case	NOT CHECKED
Customer concentration (how much revenue comes from the largest pharma advertisers)	Nobody ran it	Unmeasured. This is the actual mechanism behind the “pharma budgets are shrinking” narrative	NOT CHECKED
Detector coverage was complete	Compared both benches' checklists against the assigned set	Red ran three checks not assigned to this company (two construction-permit checks and a natural-hazard check — irrelevant to a company that leases two offices) while skipping four that were assigned	NO — red erred on coverage
Blue's “\$18.4M second-half raise”	Looked for the derivation	Not shown in the record and not reconstructible from the pack	NOT CONFIRMED

Primary sources: the 10-Q filed 2026-08-06 ([SEC filing index, accession 0001516513-26-000040](#)); the stock-compensation series above; and the company's Q1 FY27 results release

PROPOSED ENTRY BANDS

There is no entry. Both benches say stay flat, and I agree. What follows are the conditions under which we reopen — reopening means we run the case again, not that we buy.

Size today: 0% of book. No order, no starter, no resting bid. A resting order into a company with an unconfirmed earnings date would be a bet on the earnings reaction, which is not the bet we're making.

Price levels, derived from the benches' own valuation math. All use \$191.2M shares, \$687.8M net cash, \$676M revenue, \$319M reported profit, and \$147M of stock pay — so \$172M of true operating profit:

Price	What you'd be paying (after stock pay)	Meaning
\$25.89 (today)	24.8×	Where the objection lives. No action.
\$28.50	27.7×	Close the file. If it re-rates here, the market resolved the valuation argument against us and there is nothing left to court.
\$21.60	20×	Watch level. Refresh the case; still not cheap for 7% growth.
\$19.80	18×	Reopen level. Combined with the evidence conditions below, this is where a minimum starter becomes arguable. Sits just above Wells Fargo's \$18 target.
\$17.16 (52-week low)	15.1×	Add level. Worth noting plainly: on our own standard this business does not become genuinely cheap until it revisits its 52-week low.

Proposed sizing if — and only if — the reopen conditions clear: minimum starter of **0.5% of book at ≤\$19.80**, one add to a **1.5% hard cap at ≤\$17.16**. The cap is deliberately small and does not scale with conviction: with two share classes there is no acquirer bid to catch us, so the downside tail is unbounded in a way position sizing has to respect.

Evidence conditions that must clear alongside the price (all five, from the benches): 1. December-quarter revenue guide **at or above the recomputed threshold** (see the sensor note below — blue's \$193M is not yet a usable number). 2. Quarterly stock pay stays **below \$42M** (the plateau holds). 3. Deferred revenue at or above **\$109M**, *or* below it with billings still growing at least as fast as revenue (drift stable, not widening). 4. Diluted share count stays **below 195M** (buybacks still absorbing the stock pay). 5. Form 4 documents show **no unscheduled open-market founder selling** into the August bounce.

Upside reopen (the case we'd be wrong about): if the November guide comes in above the ramp threshold *and* stock pay stays flat *and* billings re-accelerate, this becomes a stabilizing business with a

shrinking share count, and we should be willing to reopen even at a higher price than \$21.60. We should not let a rejection harden into a permanent no.

Catalyst dates: - ~**2026-10-15 (estimated, unconfirmed)**: the company typically announces its earnings date about three weeks ahead. This is when the November date becomes real. - **2026-11-05 (UNCONFIRMED — derived from a data vendor, not a company announcement)**: September-quarter results and the December-quarter guide. This is the only gradeable event in the file. - **Late October**: run the case again, once the date is confirmed and the sensors below are actually armed.

Nothing between now and mid-October changes any input. There is no reason to touch this name for roughly eight weeks.

ARMED SENSORS

(armed provisionally with this deck; session adjudication ratifies)

- price alert: fires below 17.16 (deck-provisional: 52-week low, 15.1x post-stock-pay operating profit; genuinely cheap, add level to a 1.5% cap)
- filing watch: 8-K until 2027-02-17 (deck-provisional: September-quarter results and the December-quarter guide; the only gradeable event, and the guide midp)
- dated pack DOCS|2026-10-15: DOCS: expected company press release naming the September-quarter earnings date; pattern-based, roughly three weeks ahead (date confirmed=Fa
- dated pack DOCS|2026-10-26: DOCS: hard reminder to re-run this case regardless of other signals (date confirmed=False)
- dated pack DOCS|2026-11-05: DOCS: September-quarter results and December-quarter guide; date is vendor-derived, no company announcement names it (date confirmed=False)

What remains unverified

Stated plainly, without softening:

1. **The earnings date.** 2026-11-05 comes from a data vendor. No company announcement names it. Every timing statement in this deck inherits that uncertainty.
2. **Insider selling.** Six insider-related filings in two weeks, following a price bounce. Nobody opened them. We do not know the amounts or whether the sales were pre-scheduled.
3. **Customer concentration.** How much of revenue comes from the largest handful of pharmaceutical advertisers is unknown to us. This is the single largest unmeasured risk in the name.
4. **Whether the profit guidance cut was investment or deterioration.** The company raised revenue guidance by roughly \$6M while cutting profit guidance by roughly \$10M and widening the range from \$12M to \$20M. Blue argues this reads as deliberate spending on a new search product, citing a wave of San Francisco engineering job postings — but that job data is the

contaminated data. The exculpatory reading rests on evidence we've proven unreliable. **This is the one fact that could turn this from a pass into a starter, and it is unresolved.**

5. **The December-quarter seasonal split**, as above.

6. **Contracted future revenue.** The company does not disclose it. We cannot verify the backlog.

7. **The long-term portion of deferred revenue.** The billings math assumes it's negligible.

Nobody checked.

8. **The exact current share count.** We used a weighted average; the cover-page count for both share classes was never pulled. Small, but it feeds every valuation number above.

9. **Independent price verification.** Neither bench could pull independent price bars this session.

Given that both benches then quoted a price the pack does not contain, this gap is not cosmetic — the entire valuation stack rests on one unverified feed reading, which I have re-derived internally but not independently confirmed.

Honest summary of our own performance: the original brief was wrong on two of three points. The red bench correctly killed both, then broke its own price correction, misread a growth trend from two data points, ran three irrelevant checks while skipping four assigned ones, and accepted data from a tool it had just proven broken. The blue bench caught most of that, then published an alarm threshold its own arithmetic contradicts. The one genuinely new and correct finding in the entire case — that two share classes remove the takeover floor — came from red and was missed by the brief. The conclusion (no position) is right. The reasoning behind it needed two rounds of correction to get there, and the sensors that would make the conclusion durable do not exist.

```

{"price_gates": [{"level": 28.50, "direction": "above", "basis": "27.7x
post-stock-pay operating profit; the valuation argument resolved
against us, close the file"}, {"level": 21.60, "direction": "below",
"basis": "20x post-stock-pay operating profit; refresh the case,
still not cheap for 7% growth"}, {"level": 19.80, "direction":
"below", "basis": "18x post-stock-pay operating profit; reopen
level, minimum starter arguable if all five evidence conditions
clear"}, {"level": 17.16, "direction": "below", "basis": "52-week
low, 15.1x post-stock-pay operating profit; genuinely cheap, add
level to a 1.5% cap"}], "event_gates": [{"forms": ["8-K"], "why":
"September-quarter results and the December-quarter guide; the only
gradeable event, and the guide midpoint is the primary kill test"},
{"forms": ["4", "144"], "why": "six insider filings from Aug 7-17
remain unread; ongoing watch for unscheduled open-market founder
sales"}, {"forms": ["10-Q"], "why": "quarterly stock pay vs the $42M
plateau, deferred revenue current and long-term portions, and
diluted share count vs 195M"}, {"forms": ["10-K"], "why": "top-
customer concentration disclosure, the unmeasured exposure behind
the pharma-budget narrative"}], "catalyst_dates": [{"date": "2026-
10-15", "what": "expected company press release naming the
September-quarter earnings date; pattern-based, roughly three weeks
ahead", "confirmed": false}, {"date": "2026-10-26", "what": "hard
reminder to re-run this case regardless of other signals",
"confirmed": false}, {"date": "2026-11-05", "what": "September-
quarter results and December-quarter guide; date is vendor-derived,
no company announcement names it", "confirmed": false}],
"immediate_entry": null}

```

Three things worth your attention before this gets filed:

- **Both benches quoted a price the evidence pack does not contain** (\$25.21 vs the pack's \$25.89). Red used it to "supersede" the brief's valuation multiples — and the brief's original 6.3x / 13.4x / 24.9x were correct. That's a supersede that made the record worse, and it's worth a look at why both benches produced the same wrong number independently.
- **Blue's headline alarm is not arithmetically usable as written.** 53% of \$348.9M is \$184.9M, not the \$195M it stated, so the \$193M kill threshold needs the real seasonal split before it can be armed without generating false fires.
- **Zero sensors exist** for a case whose entire value is a late-October re-court. Four price alerts, an earnings-date-confirmation watcher, and the Form 4/144 fetch are all cheap and none are running.