

PITCH — SHOP — ADJUDICATED

RULING (2026-09-01, Fable full-court adjudication, principal-directed): NO POSITION — the dogfood boundary held. Honest FV \$113–140 built at the 5.2% bond floor vs tape \$141.94; fairness swings 24% on the discount rate alone = unverified fairness, which the dogfood allowance forbids regardless of how good the product is. Band $\leq$ \$110 court-on-touch (NOT pre-authorized — the frozen-software-axis add case must be argued at touch); \$94 = table-pounder review; a dogfood starter becomes available only $\leq$ \$113 with the growth path intact.

What the adjudication added beyond the benches: - **The take-rate-inversion detector shipped in net-coverage form** — and its first application is *against* the thesis: Google Universal Cart launched an owned checkout two months after OpenAI withdrew, charging merchants zero “for now.” One rail died; the net coverage of owned rails did not decline. Toll and ownership now score as separate axes. - **The AI-GMV denominator, estimated:** ~\$630M of SHOP GMV $\approx$ **0.13%** — reaching 1% ~2029, 5% ~2031-32. A 2029+ option, not a 2027 driver; sized accordingly (zero). - **The IR disclosure ask (three metrics):** GMV by referral origin with AI broken out + definition; Shop Pay attach on agent-referred vs all sessions; GMV via third-party agentic rails completing off Shop Pay. The first and third grade our frozen 0.20 call; the third is the one they’ll resist. - **The CSP tripwire** (principal’s puts question): strike $\leq$ \$110, IV>RV, premium $\geq$ 8% annualized after the ~50% tax haircut. Fails hard today (IV 40 vs realized 72, 0th percentile — selling insurance below cost); the sanctioned window is the ~Nov 3 pre-print IV ramp. - **The dogfood datum stays an observation:** the only public generalization evidence (Shop Pay’s conversion-lift study) failed provenance — sponsor-commissioned, unnamed firm, “up to” framing. Observer row 14 (your monthly SHOP-vs-AMZN split) remains the live instrument. - **The FV-unlock is named:** two consecutive quarters of gross-profit dollars growing $\geq$ revenue moves the band UP — management currently guides ~500bp the other way, verbatim. Frozen at 0.15 for Q3.

The bench pitch below is preserved as written (pre-adjudication).

What we believe

Shopify is an excellent business that is currently accelerating, and we do not think we should buy it today.

The business facts are not in dispute between our two benches. In the June 2026 quarter: merchandise volume across the platform was \$115.6bn, revenue was \$3.58bn (up 33.7% from a year earlier), operating income was up 67.7%, free cash flow was \$654M, the company holds

\$4.95bn in cash and securities with no disclosed debt, and the diluted share count is *shrinking* — 1,297.9M versus 1,309.0M a year ago. Five straight quarters of 30%+ growth. That is a rare combination and we should say so plainly before we say anything negative.

What we believe that is *ours*, not borrowed:

1. **The “AI killed Shopify’s checkout” story is backwards on the merchant P&L.** When OpenAI ran Instant Checkout inside ChatGPT, merchants paid OpenAI a 4% fee **on top of** Shopify’s own fees. OpenAI retired that rail in March 2026. The withdrawal deleted a 4-point toll stacked above Shopify’s rail and pushed the same volume back through Shopify Payments. Most commentary scored the withdrawal as Shopify losing a channel. On the merchant’s income statement it was a margin restoration.
2. **The AI-shopping channel is too small to matter for years.** Our own bottom-up estimate: roughly 900M weekly ChatGPT users × ~5% with shopping intent in a given week × ~30% clicking out × ~3% converting × ~\$120 average order ≈ **\$2.5bn a year of AI-originated volume across the entire internet**. Give Shopify a generous 25% of that and it is ~\$630M a year — about **0.13%** of Shopify’s ~\$490bn annual volume. Even doubling every year, it reaches 1% of volume around 2029. It contributes under half a percentage point to 2027 revenue growth. This is the single most useful number in the whole court, and it argues the AI angle is an *instrument to watch*, not a reason to pay up.
3. **At \$144 we are paying above our own honest estimate of what it’s worth.** Our discounted-cash-flow, built at the discount rate our own doctrine mandates (5.2% long yield + 4.5% equity risk premium = 9.7%), gets to **\$139.90** — and that path already assumes revenue compounds 4.1x over ten years *and* free-cash-flow margin doubles from 18% to 28%. At the more honest 11% rate for a stock whose own 52-week range spans nearly 2x, fair value is **\$113**. Today’s price sits above the top of that band.

What the market believes

The market believes Shopify is a durable, category-winning commerce platform whose growth decelerates gracefully — and it is willing to pay roughly 59x next year’s earnings and ~68x current-year free cash flow for that.

Concretely: 52–53 analysts cover it, essentially no one recommends selling it (one “strong sell” out of 53), and the average price target is \$167–\$171 versus a \$144 price. The load-bearing consensus number is 2027 revenue of about **\$18.7–19.1bn, i.e. +24% to +26% growth**, with earnings of \$2.45 a share.

So the market’s view is: growth steps down from 33% to ~25%, margins keep improving, and the multiple holds. To get to a \$170 target you have to either use a cost of equity near 8.5% or assume a terminal multiple above 20x cash flow. That is not a crazy view. It is a *rate-sensitive* view.

Why we might have edge

Three candidates, ranked by how much we actually trust them.

1. The toll-stack insight (real, verifiable, and we found it). The 4% ChatGPT-checkout fee is documented, and its removal is documented. The general lesson is worth more than the Shopify application: *when a big platform launches and then withdraws a checkout rail, the fee it stacked on top of the incumbent's take rate disappears too — so an apparent loss of distribution can be a margin gain.* We have written this up as a mechanism for the knowledge graph. That is a durable asset regardless of whether we ever own the stock.

2. The growth-deceleration gap (small, gradeable, unfunded). If Shopify holds ~30% growth into 2027 against the street's 25.6%, revenue lands near \$19.8bn and earnings near \$2.70 — roughly **+10% versus consensus at an unchanged multiple**. That is a genuine, measurable divergence. It is also only 10%, which does not clear the bar for entering at 59x. And critically, we have no *mechanism* to underwrite the beat: the channel we'd point to (agentic commerce) has no disclosed denominator anywhere in the filings.

3. The product observation (honest but unsupported). The original thesis rested partly on direct experience that Shopify's checkout converts better. It may well be true. But the only quantitative support anyone has is a study Shopify itself commissioned in **April 2023** with an unnamed "Big Three" consultancy, methodology unpublished, headline framed as "up to 50%" and "up to 36%" — a construction that cannot be proven wrong. We should treat that as marketing, not evidence, even when we suspect the conclusion is right.

Why it might be priced in

Bluntly: our growth path is *above* consensus by 1.7%, which is nothing. We are not smarter than the street on the numerator. Where we differ is the discount rate — and our difference points **down**, not up. Buying a stock because we agree with consensus growth while disagreeing with consensus discounting is a bet on interest rates dressed up as a stock idea.

Two further reasons to think it's priced:

- **Coverage and crowding.** 52–53 analysts, \$186bn market value, price in the middle of its own range (21% below the high, 53% above the low). This is not an under-followed situation. Under our own rules, crowding conditions *size and timing*, it is not itself a reason to reject — but it does mean there is no discovery gap to harvest.
- **Management is guiding margin down, and our fair value needs margin up.** The Q3 guide is revenue growth in the "low thirties" but gross profit dollars only "mid-to-high twenties" — roughly 500 basis points of mix dilution the company is telling us about in advance. Gross margin was 47.7% this quarter versus about 48.7% a year ago. Payments and other lower-margin merchant services are now 77.6% of revenue. Our \$139.90 fair value requires cash-flow margin to *double*. The company is guiding the underlying margin *down*.

What could go wrong

(Pre-mortem. Written as if we had bought at \$144 and lost money.)

Failure narrative 1 — The rate re-rate. (Monitorable.) Nothing about the business breaks. Growth comes in at 28%, respectable. But long yields move up 100bps, the market applies our 11% discount rate instead of its 8.5% one, and the multiple goes from 59x to 40x. The stock

trades \$95–\$110 with earnings *beating*. We lose 25–30% on a thesis that was operationally correct. **Tripwire:** 10-year yield, and the gap between reported growth and the multiple. This is the most likely way we lose, and it is fully watchable.

Failure narrative 2 — Net rail coverage rises against us. (Partially monitorable — the tripwire exists but the measurement does not.) Our own red bench caught this and it is the sharpest thing in the court. OpenAI retiring Instant Checkout was the bullish data point. But **Google launched Universal Cart at I/O in May 2026, with Google Pay checkout live on Shopify merchants including Fenty and Steve Madden** — on the platform that actually owns shopping-intent search volume. Three months after the “bullish” event, a *bigger* rail landed on Shopify’s own merchants. Net third-party checkout coverage went **up**, not down. If the toll-stack story is right on OpenAI but wrong in aggregate, Shopify’s payments attach rate erodes quietly and the margin path never arrives. **Why the tripwire is weak:** Shopify does not disclose volume by referral origin, does not disclose Shop Pay attach on agent-referred sessions, and does not disclose volume transacted on third-party rails. We would learn this from the margin line eighteen months late.

Failure narrative 3 — No tripwire exists. (Unknown-unknown class: *unpriced-by-construction* — risks where no reported line item can move first.) Two live shapes here, neither of which any ratio we track would signal in advance: - **Governance concentration.** The share structure includes a founder share class (visible in the XBRL as `shop:FounderShareMember`). A key-person discontinuity at Shopify is a step function, not a trend. No metric decays first. - **Category redefinition.** The thing that ends “merchant owns the storefront” is not slower growth in storefronts — it is a world where the merchant-of-record layer itself collapses into a model or a marketplace, and the number of storefronts stays flat right up until it doesn’t. Our entire sensor set measures *volume through the current architecture*. It is structurally blind to the architecture changing.

We name this class explicitly because our instinct is to answer every risk with a monitor, and here we honestly cannot. The correct response to an unmonitorable risk is *smaller size*, not a fake sensor.

The reflexive and structural leg. - Our size versus liquidity: not a constraint. Shopify is a \$186bn company; our contemplated position (0.5–1.5% of book) is a rounding error against daily volume — we estimate well under 0.01% of average daily volume, though we did not verify an ADV figure this session and are stating it as an estimate, not a measurement. We cannot move this stock and we cannot get trapped in it. Liquidity is the one risk this name doesn’t carry. - **Predicted flow — this is the real structural trap.** Our proposed reopen gate is \$110. Ask honestly *what brings the price to \$110*. Almost certainly: a quarter where growth breaks below 25%, or gross profit dollars decelerate again, or the rate environment shifts. In other words, **the price we say we want arrives attached to the news that invalidates the thesis**. Our own value-ladder rule warns against bidding into flow we ourselves predicted. A resting \$110 order with no accompanying condition on *why* it got there is an accidental bet that the de-rate is technical rather than fundamental. The reopen band below is therefore written as **price AND unchanged growth path**, never price alone. - **Doctrine interactions pulling us toward a mistake.** Three of our own standing rules push toward buying here and we should name them so they don’t operate on us invisibly: (a) the *dogfood pay-up allowance*, which lets us pay modestly rich for a product we use and believe in — but it explicitly requires fairness to be **verified**, and our

fair value swings 24% (\$113 to \$140) on the discount rate alone, so fairness is *not* verified and the allowance does not apply; (b) the *clean-court minimum-starter default*, which says a court with no named kill should produce a starter — the blue bench supplied a named kill (undisclosed agentic denominator + a 10% edge against a 59x multiple), so the default is properly overridden rather than ignored; (c) the *don't-default-to-wait* rule on high-quality names, which is designed to stop us from admiring businesses into oblivion — it applies at *fair*, and we are above our own fair band.

If this position loses money, the most likely reason will be ____

____ = that we paid up for a business we were correctly right about, and the multiple — not the business — did the losing: growth landed near consensus, the discount rate went up, and 59x became 40x while earnings beat.

What we checked ourselves

Claim	How checked	Found	Confirmed?
Price is above our honest fair value	Reverse-DCF at doctrine 9.7% and at 11%, on audited cap structure	\$139.90 / \$113 versus \$144 tape	Yes — but note the 24% swing on discount rate alone is itself the finding
We have a growth edge vs. the street	Our bull path vs. 52-analyst consensus	Our path \$19.0bn vs street \$18.7bn = +1.7%. Essentially no edge on growth	Yes (edge does <i>not</i> exist)
“Forward P/E ~78x — fully priced” (our own brief)	Recomputed on FY26 and FY27 consensus with Q2’26 diluted shares	75.5x FY26, 58.9x FY27 . Our own brief froze the wrong year. Still expensive; not 78x	Bench error — corrected
“45 analysts, 34 Buy / 10 Hold / zero Sell , target \$170.72” (our own brief)	Read the coverage page directly	53 analysts, 30 Strong Buy / 11 Buy / 11 Hold / 1 Strong Sell , target ~\$171	Bench error — the “zero sells” rhetorical hinge was factually wrong
Our two benches agree on the consensus numbers	Side-by-side	They don’t : red read 52 analysts / PT \$167 / FY27 \$18.689bn; blue read 53 / \$171.15 / \$19.13B. Same source, different reads	No — unresolved. Do not treat either as precise; treat “~53 analysts, ~\$170, ~+25%” as the honest resolution
~70x free cash flow	18% guided margin × FY26 consensus revenue	68.4x	Yes
Instant Checkout was withdrawn	Two independent reports	Yes — though our benches disagree on the date (one says 2026-03-04, one 2026-03-20). Fact holds; date is unresolved	Fact yes, date no
The withdrawal was bearish for Shopify	Traced merchant economics	Overtured. Merchants paid OpenAI 4% <i>on top</i>	No — our original brief conceded this away in

Claim	How checked	Found	Confirmed?
		of Shopify's fees from 2026-01-26; removal restores 400bp and routes volume back to Shopify Payments	a parenthetical and was wrong to
"Rail withdrawal = bullish" is a general detector	Looked for the disconfirming case	Refuted as written. Google launched Universal Cart (May 2026) with checkout live on Shopify merchants. Net rail coverage rose. Detector was built on n=1 with its own inversion unlogged	Bench error — the thesis selected the confirming row. Detector must be rebuilt as <i>net</i> coverage across all referring platforms
The Shop Pay conversion-lift evidence is solid	Provenance check on the source	Sponsor-commissioned, April 2023, consultancy unnamed, methodology unpublished, "up to 50%" framing	No — does not survive provenance. It was our only quantitative support
AI-originated volume is material by 2027	Bottom-up denominator built from scratch	~\$630M/yr to Shopify $\approx$ 0.13% of volume; reaches 1% ~2029	Yes (it is immaterial)
Merchant satisfaction is holding	App-review velocity detector, executed	Merchant admin app: recent cohort 2.22 vs lifetime 4.69, delta -2.47 — negative . But only 1.07 reviews/day, and written reviewers skew angry	Weak signal only. Too thin to carry a churn conclusion on a multi-million-merchant base
Consumer preference for merchant-owned checkout	Tried to resolve the consumer Shop app	UNRESOLVED — could not obtain a store ID. Zero independent evidence on <i>either</i> side	No — reportable blind spot
The 2026-08-31 13D/A is insider selling	Read the filing	It is Shopify filing on Klaviyo , not Shopify insiders. Stake ~\$0.5bn $\approx$ 0.3% of market cap — immaterial	No — and catching this prevented a false kill
Dilution or debt risk	Share series, SBC, balance sheet	Shares falling 0.84% YoY, SBC ~\$128–150M/qtr (~3.4% of revenue), no debt	No risk found — declined to manufacture one
Shares are falling because of buybacks	Searched for a repurchase authorization	Could not find one. May instead be a mechanical effect of the lower share price on dilution math	No — plausible, not confirmed. Open item

Claim	How checked	Found	Confirmed?
Committed backlog / revenue visibility	RPO drift detector	Uncheckable — Shopify discloses no RPO. Also: disclosed MRR of \$221M annualizes to ~\$2.65bn against a ~\$3.2bn subscription run-rate; the gap is unexplained	No — genuine coverage gap
We have a live price	IBKR gateway	Permission denied to both benches. No bench in this court saw a fresh quote; we used the \$144.17 pack tape	No — stated, not skipped
The evidence pack was current	Cross-checked against live filings	Pack XBRL stopped at Q2- 2025 ; pack volume figure (“~\$80-90bn quarterly”) was stale against actual \$115.6bn. Both benches independently caught and superseded it	Pack defect — fixed by both benches
Red case was properly adversarially audited	Record inspection	The blue bench’s “red case you are auditing” block arrived empty ; it audited the brief instead. This court’s adversarial loop was structurally incomplete	No — process defect, disclosed

Primary source of record for the quarter, quoted over any transcript or secondary report:
<https://www.sec.gov/Archives/edgar/data/1594805/000159480526000046/exhibit991pressreleaseq220.htm>

PROPOSED ENTRY BANDS

Both benches say NO ENTRY at \$144. There is no immediate order. What follows are reopen conditions, taken from the benches’ own math.

Reopen band 1 — \$110.00 or below, WITH the growth path intact. \$110 is below the \$113 fair value computed at the honest 11% discount rate. “Intact” is a hard condition, not a formality: quarterly volume growth still $\geq 25\%$ and revenue growth still $\geq 25\%$ as of the most recent print. This condition exists specifically because the flow that takes the stock to \$110 is most likely a growth disappointment — if it *is* a growth disappointment, the gate does not fire and we re-court from scratch. → Size on fire: **0.5% of book**, entered in two tranches (crowded names get calendar-spaced entries, never one clip).

Reopen band 2 — \$94.00 or below (the 52-week low), same growth condition. The red bench called this a table-pounder and we agree. At \$94 the market is applying roughly a 12–13% discount rate to a business compounding above 30%. → Size on fire: scale to **1.5% of book** total.

Reopen band 3 — evidence-based, price-independent: gross profit dollars grow at or above revenue for two consecutive quarters. This directly refutes the margin-compression objection and unlocks the cash-flow-margin expansion our \$139.90 fair value requires. If this fires while the stock is between \$110 and \$144, we re-court immediately rather than waiting for a price gate. Earliest possible confirmation: the Q3 print, with Q4 as the second leg — i.e. not before roughly February 2027.

No exits or trims specified — we own nothing.

Catalyst dates. The only one is **2026-11-03, the expected Q3 print — and it is UNCONFIRMED.** It is derived from a data vendor, not from a Shopify press release. Approximately 44 trading days out as of this writing, well outside any print-proximity window, so no print-decisive work is owed yet. Someone must confirm the date against a company announcement before we treat it as a real calendar item.

Two things that would change the shape of this pitch entirely, not just the price: - A named, published, third-party (non-sponsor) checkout-conversion study with disclosed methodology. That would restore the product leg we currently cannot support. - Google materially retrenching Universal Cart checkout, or Shopify disclosing that agent-referred sessions convert on Shop Pay at parity or better.

The disclosure ask (what we would tell Investor Relations we need): (a) volume by referral origin, with AI/agentic surfaces broken out as their own line, quarterly, with “AI-originated” defined; (b) Shop Pay attach rate on agent-referred sessions versus all sessions; (c) volume transacted via third-party agentic rails where checkout happens off Shop Pay. **(c) is the one that matters and the one they will resist** — it is the direct measurement of the Google-rail risk. With (a) and (c) the thesis becomes underwritable. Without them it is unfalsifiable, and unfalsifiable things cannot be sized.

ARMED SENSORS

(armed provisionally with this deck; session adjudication ratifies)

- price alert: fires below 94.0 (deck-provisional: 52-week low; implies a ~12-13pct discount rate on a business compounding above 30pct - scale starter t)
- filing watch: 8-K, 10-Q until 2027-02-28 (deck-provisional: First disclosure of GMV by referral origin, Shop Pay attach on agent-referred sessions, or GMV on thir)
- dated pack SHOP|2026-11-03: Q3 print (UNCONFIRMED) - OBSERVATION: a disclosed AI/agentic GMV denominator triggers the re-court; GMV growth <+25pct re-examines damage-ab

What remains unverified

Stated plainly, because these are the reasons the conviction here is capped:

1. **No live price.** IBKR was permission-denied to both benches. Every number in this deck keys off the \$144.17 pack tape. Nobody in this court saw a fresh quote.
2. **Consensus figures disagree between our own benches** — 52 vs 53 analysts, \$167 vs \$171 target, \$18.689bn vs \$19.13bn 2027 revenue, from the same source. We have not reconciled them. Treat consensus as “~+25% growth, ~\$170 target” and no more precise than that.
3. **The Instant Checkout withdrawal date** is reported as both 2026-03-04 and 2026-03-20 by our own two benches. The fact holds; the date does not.
4. **Whether the shrinking share count is a buyback.** No repurchase authorization located. Could be a treasury-stock-method artifact of the lower price.
5. **Consumer-side evidence: entirely absent.** The consumer Shop app could not be resolved. The merchant-side signal that *did* come back (recent-cohort rating 2.22 vs 4.69 lifetime) is genuinely negative but rests on 1.07 reviews per day — too thin to conclude anything, and written reviewers skew angry.
6. **No revenue-visibility measure exists.** Shopify discloses no committed backlog. Given that the entire argument is about whether 2027 growth is 25% or 30%, this is a material gap and there is no workaround.
7. **The size of agentic commerce is our estimate, not anyone’s disclosure.** The 0.13%-of-volume figure is a five-assumption chain we built ourselves. It is directionally robust — it would take an order of magnitude of error to change the conclusion — but it is not measured.
8. **This court’s adversarial loop did not fully close.** The blue bench received an empty red-case record and audited the original brief instead. The red bench’s findings were never independently rebutted. That is a process defect and it caps how much weight adjudication should give this deck.
9. **The 2026-11-03 print date is unconfirmed.**

Honest summary of where the two benches actually landed: red rejects at 8/10 conviction, explicitly labelling it a *valuation and evidence-provenance* kill rather than a business kill, and says a \$110 print reverses them quickly. Blue overturns much of red’s *reasoning* — the multiple axis, the coverage count, the “unmeasurable” concession — and still lands on flat-no-starter, at 7/10, because a ~10% edge does not clear a 59x entry. Two benches, different roads, same destination. That convergence is the most useful thing this court produced, and it is worth more than either bench’s individual argument.

```
{"price_gates": [{"level": 110.0, "direction": "below", "basis": "Below the $113 DCF fair value at an honest 11% discount rate; arms ONLY jointly with GMV and revenue growth still >=25pct, because the flow that produces $110 is most likely the growth break that invalidates the thesis"}, {"level": 94.0, "direction": "below", "basis": "52-week low; implies a ~12-13pct discount rate on a business compounding above 30pct – scale starter to 1.5pct of book, same joint growth condition"}], "event_gates": [{"forms": ["8-K", "10-Q"], "why": "First disclosure of GMV by referral origin, Shop Pay attach on agent-referred sessions, or GMV on third-party agentic rails – any one re-opens the court regardless of price"}, {"forms": ["8-K", "10-Q"], "why": "Locate a share repurchase authorization; the 0.84pct YoY share-count decline is currently unexplained and may be a treasury-stock-method artifact of the lower price"}, {"forms": ["8-K"], "why": "Quarterly gross profit dollar growth at or above revenue growth – two consecutive quarters refutes the guided ~500bp mix dilution and unlocks the FCF-margin path the $139.90 FV requires"}], "catalyst_dates": [{"date": "2026-11-03", "what": "Expected Q3 2026 print – vendor-derived, must be confirmed against a company announcement; disclosed AI/agentic GMV denominator triggers re-court, GMV growth below +25pct re-examines the damage case", "confirmed": false}], "immediate_entry": null}
```