

PITCH — TTAN — PROPOSED — PENDING ADJUDICATION

ADJUDICATED 2026-09-16 — STARTER 0.4x — stage after the 9/21 flow gate (conviction 7/10).

What we believe

ServiceTitan sells software to home-services contractors (HVAC, plumbing, electrical). It trades at \$56.285, down 52.9% from its 52-week high and only 3.9% above its 52-week low, after the September 8 print. Our belief, stated plainly:

1. **The subscription business is not broken.** Subscription revenue was \$212.4M in the July quarter, growing ~22%. Contracted revenue due within the next twelve months (cRPO) grew roughly +23% year over year — faster than reported revenue (+21%) and much faster than the guided next quarter (+14.8%). A business whose customers were actually leaving cannot have a twelve-month forward book growing 23%.
2. **The damage is real but localized.** It sits in the transaction-linked “usage” line (payments/GTV) and in a deliberately shrinking professional-services line that carries a **-158% gross margin** — i.e., the company loses money on it on purpose and is winding it down (10-Q, 2026-09-08).
3. **The price now reflects a much worse business than the one in the filings.** At ~95.4M shares and \$56.285, market cap is ~\$5.37B against implied FY27 revenue of ~\$1.14B — **4.70x sales** for a company still growing ~18% for the full year with positive free cash flow (~\$50.5M in the quarter, ~\$56.4M six-month operating cash flow).

So: **reduced damage-arriving**, not “nothing happened” and not “the business is failing.”

What the market believes

The market believes growth is stepping down permanently and that management just told you so. Its case is not stupid:

- Guided next-quarter growth is **+14.8%**, versus **+25.0%** in the same quarter a year ago. That is roughly 1,000 basis points of deceleration in four quarters, and the comparison was *not* harder — last year’s Q3 actually grew faster than last year’s Q2.
- The implied fourth quarter (~\$293.9M, +2.8% sequential, ~+13.9% year over year) means the full-year guide **back-loads no recovery at all**. Nobody is being asked to wait for a reacceleration; there isn’t one in the numbers.
- Gross transaction volume grew +17% — slower than the software — and management did not explain the gap.

- The company simultaneously announced a strategic pivot toward its AI product (“Max”) over broader trade expansion, on the same day as a chief revenue officer change. The market read that as “the old growth engine is being replaced,” which is a re-rating event, not a quarter event.
- Retention is disclosed as a **band** — “greater than 110%” — not a number. A move from 118% to 110.5% would look identical in that disclosure ([8-K EX-99.1](#)).

Why we might have edge

- **We decomposed the guide correctly and the loud bear case did not.** The headline bear argument compares a *conservative forecast* to a *reported result*, which manufactures a fake decline equal to the company’s habitual beat. Guide-to-guide is \$285M → \$286M, **+0.35%**. Apply the demonstrated beat (+2.74% last quarter, +3.7% on the year-ago Q1 guide) and next quarter lands near **\$293.8M, +0.35% sequentially** — a modest seasonal break versus the prior year’s +2.9%, not a collapse.
- **We separated the two things inside the “non-subscription” bucket.** Strip out professional services (shrinking on purpose, negative margin) and the implied usage line is about **-9.0% sequentially**, or **-7.4%** after adjusting for the one fewer business day management named. That is a genuine dent. It is not the -16% to -22% the bear decomposition produced.
- **We ran the bear’s own falsification test.** The bear said “if cRPO grows $\geq 22\%$, my backlog argument collapses.” Using the company’s disclosed RPO and its next-twelve-month percentage, cRPO is ~\$263.7M versus ~\$214.2M, **+23.1%**. The test was run and it failed for the bear. (Caveat below — the percentages are rounded, and that matters.)
- **We are not carrying the position.** Book is flat, no orders. We can choose our entry rather than defend one.

Why it might be priced in

Symmetrically, here is the honest list of reasons this is *not* edge:

- **4.70x sales is not obviously cheap for a ~14% exit-rate grower.** The “cheap” frame is priced off a 21% grower. If the run-rate is 14%, 4.7x is roughly fair, not a gift. This is the single strongest argument against us and we do not have a rebuttal beyond “the forward book says 23%.”
- **Duration is quietly extending.** The share of the backlog due within twelve months slid **53% → 52% → 50%** across three consecutive quarters. Longer contracts bought with price concessions look like backlog strength and aren’t. It’s a ~300 basis point effect, not a 1,570 basis point illusion — but it is real and it moves the same direction three quarters running.
- **Dilution is genuine.** Diluted shares rose 1.50% in a single quarter. Even discounting the post-IPO normalization distortion, the share count is a headwind nobody is being paid for.
- **The AI-pivot story is management-declared, not inferred.** That makes it the most credible bear leg on the table, and it is precisely the kind of thing that takes a multiple down and keeps it down for three or four quarters.

- **The tape has not absorbed anything.** Pinned 3.9% off the low with a -52.9% drawdown is not a stabilized dislocation. Neither of our benches could pull a single price bar to measure absorption; we are eyeballing this.
- **We cannot see the street's number.** No consensus pull succeeded on either bench. "The market overreacted" is being asserted against an anchor we cannot see.

What could go wrong

Failure narrative #1 — The usage line is the whole story and it keeps going (most likely).

The implied ~7.4% sequential usage decline, adjusted for calendar, is not phasing; it is contractors doing less volume, or a lower take on that volume. It compounds for two more quarters, GTV growth falls under 12%, and the software line follows because usage and seats are correlated at the same customer. *Monitorable*: yes — usage revenue and GTV disclosed at the 2026-12-03 print; Austin and Bernalillo County trade-permit counts are readable monthly and independently of the company. **Both permit checks are currently unrun. That is our own gap, not the company's.**

Failure narrative #2 — The AI pivot is a repositioning that eats two years. Management deprioritizes trade expansion for the Max product, sales leadership turns over, the new motion takes four to six quarters to produce bookings, and the multiple compresses to 3.5x while we wait. *Monitorable*: partly — cRPO and net dollar retention at each print, plus sales-leadership Form 4s and 8-K personnel items. But if the market simply decides to pay less for the same numbers, no tripwire fires.

Failure narrative #3 — THE ONE WITH NO TRIPWIRE: the payment processor gets repriced.

The 10-Q carries a third-party payment-processor concentration disclosure sitting directly under the usage line. If that processor renegotiates the economics — a take-rate split change, a term change, a migration — revenue on that line falls and **we will learn about it only when it prints.** There is no filing, no permit series, no satellite, no app store, no alt-data vendor that observes a private bilateral contract renegotiation. **Unknown-unknown class: PRIVATE BILATERAL COUNTERPARTY-TERMS RISK — economics governed by a confidential commercial contract between two private counterparties, with no public observable and no leading indicator, disclosed only ex post through the revenue line it damages.** We are naming it rather than pretending a sensor covers it. A second class sits alongside it: **BANDED-METRIC MASKING** — retention is disclosed as ">110%," so deterioration from 118% to 111% is invisible by construction, and no monitor can see inside the band.

The reflexive / structural leg. Our own size is not the problem. At an assumed \$250M book, the full 0.60% authorization is ~\$1.5M, roughly 26,600 shares — against an estimated 2.5–4.0M share average daily volume, that is **0.7% to 1.1% of a day.** We move nothing and we can exit in an hour. (Volume is an estimate; neither bench pulled a bar. It is item 5 in the unverified list.) The real structural risks are two:

- **Predicted flow against us.** Four Rule 144 notices in eleven days plus three Form 4s in August 2026 is continuous insider supply. We do not know the share counts or whether they are pre-

set 10b5-1 sales — the “they knew something” version is unverified and we are not asserting it. What we *are* asserting is that there is a seller in the tape and we would be adding into it.

- **Crowded gates.** Our four go/no-go conditions at the December print (cRPO, GTV, usage sequential, retention band) are the same four every analyst will publish. If they fail, everyone exits through the same door at the same second and our “exit, don’t add” rule executes 15% below where we decided to use it. Our staging rule also makes us a **momentum buyer by construction** — we only add on confirmation, which means we pay up after the news is out and we are structurally guaranteed to have our smallest size at the low.
- **Doctrine interaction.** We are leaning on the “guide-versus-actual artifact” pattern (bear decompositions that treat a conservative forecast as a real sequential decline). That doctrine is correct in general and is exactly the doctrine that will talk us out of a genuine demand break the one time it is genuine. It cuts both ways and we should expect it to cost us here at some point.

****If this position loses money, the most likely reason will be ____**** — the usage/GTV line kept deteriorating for reasons governed by a private processor contract and by contractor demand we never instrumented, while the market re-rated a ~14% grower down toward 3.5x sales and our confirmation-only staging rule left us buying strength that reversed.

What we checked ourselves

Claim	How checked	Found	Confirmed?
Tape basis \$56.285, +3.9% off low, -52.9% drawdown	Recomputed from evidence-pack tape fields	$(56.285-54.17)/54.17 = +3.90\%$; $56.285/119.6-1 = -52.94\%$ — matches pack exactly	CONFIRMED (pack)
OUR OWN BENCH ERROR #1: red bench used px \$57.30 , blue bench used px \$56.625 ; blue then accused red of quoting an “unsourced, not-in-the-pack” price	Compared both to the shared pack	Both are wrong. Pack says 56.285. Blue’s own price is equally unsourced and its “pct_off_low 4.5%” is wrong (pack: 3.9%). Every valuation number on both benches inherits the error	BOTH BENCHES ERRED — corrected here to 4.70x, not red’s 4.79x or blue’s 4.73x
OUR OWN BENCH ERROR #2: red declared “no network — web_search unavailable” and attacked only on arithmetic	Blue bench ran the same searches and SEC fetches successfully	Red’s coverage gap was self-inflicted ; five of red’s twelve findings leaned on it	RED ERRED
Next-quarter guide implies usage revenue -16% to -22% sequentially	Rebuilt the arithmetic both ways	Guide-vs-actual gives -16/-22%. Guide-vs-guide is +0.35% (\$285M→\$286M); beat-adjusted at the demonstrated +2.74% gives ~\$293.8M and a residual of -9.3% ; removing the shrinking professional-services line leaves usage at -9.0% , or -7.4% after the one fewer business day	RED’S MAGNITUDE REFUTED; the damage itself CONFIRMED
Deceleration from 25% to ~15%	Recomputed from filings	Q3 FY26 actual \$249.163M vs \$199.275M = +25.0% (<u>10-Q</u>); Q3	CONFIRMED — this is the strongest bear fact

Claim	How checked	Found	Confirmed?
		FY27 guided +14.8% ; Q1 FY27 \$268.824M, +24.6% (10-Q); implied Q4 FY27 ~\$293.9M = +13.9% YoY , no back-loaded recovery	and it survives everything
cRPO +23.1% year over year clears the bear's own $\geq 22\%$ falsification trigger	Recomputed, then stress-tested the rounding	Point estimate: \$527.4M $\times 50\%$ = \$263.7M vs \$404.1M $\times 53\%$ = \$214.2M = +23.1% . But "50%" and "53%" are rounded. The true range on $\pm 0.5\text{pt}$ rounding is +20.8% to +25.3% (IR supplemental)	PARTIALLY CONFIRMED — BLUE OVERCLAIMED. The overturn clears the 22% bar only if rounding breaks our way. We are treating it as "likely but not settled," not as a decided kill
"The company chose a weaker retention disclosure and that choice is information"	Checked disclosure history back to the S-1	The ">110%" band is the unchanged convention since the IPO . No change of form occurred	RED'S INFERENCE OVERTURNED — but the resulting blind spot is real and we carry it as a risk
"\$72.0M GAAP-to-non-GAAP wedge is stock comp = 24.6% of revenue"	Compared against pack XBRL	Pack stock comp is \$43.7M (20.3% of revenue) . The wedge also contains intangible amortization, employer taxes, and a co-founder RSU charge visible in the 10-Q tags. "Free cash flow net of stock comp is -\$15M to -\$22M" is overstated	RED'S MAGNITUDE REFUTED; direction sustained
"+6.1% annualized dilution"	Checked the base period	Annualizes a first-post-IPO-year step (90.334M $\rightarrow$ 91.688M) when the prior pack count is 35.09M — an IPO normalization artifact. Quarterly step of +1.50% is real; the annualization is not	RED'S MAGNITUDE REFUTED
Share count / valuation denominator	Attempted to source 95.4M	Pack's newest diluted count is 2025-07-31 — twelve months stale , and the 10-Q shows a Class C line at 2026-08-31 that neither bench counted. 95.4M is sourced nowhere . Both benches' per-share math is unverified, including ours	UNRESOLVED — flagged, and item 2 below
App-store rating collapse as evidence of user churn	Checked the sample size	~6 reviews in 90 days; lifetime rating 2.95 means the app was always poorly rated, so there is no change to infer	RULED NOISE — red flagged it honestly and declined to lean on it; we agree
Payment-processor concentration is a new, undisclosed dependency	Checked the tag's period	The tag appears for the prior-year period too — recurring disclosure, not a new event	RED'S FRAMING REDUCED; the exposure itself is real and untrippable (see pre-mortem #3)
OUR OWN BENCH ERROR #3: red's detector list	Cross-checked against dispatched detectors	Red omitted three dispatched entries (clinicaltrials_lookup, common_control_merger_accounting, competitor_trial_omission). All three come back clean, so no finding changed — but the omission is a process defect	RED ERRED (procedural)

Claim	How checked	Found	Confirmed?
"One fewer business day" explains the guide	Sized it	One day of a \$72.1M quarterly usage line ≈ \$1.1M — about 15% of the gap, not all of it. Also note: this is sourced to a call transcript , a secondary source. Not contradicted by the 10-Q, but weaker than a filing	CONFIRMED as partial, REJECTED as the full explanation
Full-year FY26 base used for all year-over-year framing	Traced it	~\$965M is derived , not fetched. Every Q4 year-over-year frame on both benches carries about ±1 point	DERIVED — stated, not hidden

PROPOSED ENTRY BANDS

Total authorized: 0.60% of book (= 0.4× a 1.5% house full-size unit; the 0.4× multiplier is the blue bench's, the 1.5% unit is desk convention and is not from the benches). **No immediate fill**. Three tranches of 0.20% each. Multiples below use implied FY27 revenue of \$1,141.5M and an assumed 95.4M shares — both flagged as unverified.

Tranche	Size	Trigger	Level / condition
T1 — starter	0.20% of book	Price only, limit orders, no market orders	\$52.00–\$54.50 (4.35x–4.55x). This is at and just under the 52-week low of \$54.17 — we are deliberately bidding into the knife rather than chasing it
T2 — evidence add	0.20% of book	2026-12-03 print (UNCONFIRMED date), all four gates pass	Gates: cRPO ≥20% YoY; GTV ≥15%; usage revenue ≥ flat sequentially; retention still >110% . Pay up to \$70.00 (5.85x). Any <i>two</i> gates failing → exit the whole position, do not add
T3 — confirmation add	0.20% of book	Second confirming print (Q4 FY27, ~March 2027, unconfirmed) with usage revenue positive sequentially	Pay up to \$78.00 (6.5x)

Exits and trims - Hard stop: two consecutive weekly closes below **\$47.86** (4.0x). Below 4x the market is pricing a sub-12% grower and our thesis is simply wrong. - **Gate stop:** two of four December gates fail → full exit at the open, regardless of price. Accept the slippage; the crowding risk is named in the pre-mortem. - **Trim:** sell half at **\$71.79** (6.0x) if usage revenue is still negative sequentially at that point — that is a multiple re-rate we did not earn. - **No adds between \$54.50 and the print.** The middle is where we would be paying for hope.

Catalyst dates - 2026-12-03 — Q3 FY27 print. UNCONFIRMED (derived from a data vendor; no company press release names it). ~58 sessions out. This is a reconstruction after a print, not a race into one. - **~2026-12-04/05 — 10-Q filing**, which carries the RPO footnote and the concentration language. - **~March 2027 — Q4/full-year print.** Unconfirmed. - **Ongoing — Rule 144 / Form 4 cadence.** Four 144s in eleven days in August; the next cluster tells us whether supply persists.

ARMED SENSORS

(armed provisionally with this deck; session adjudication ratifies)

- price alert: fires below 47.86 (deck-provisional: hard stop at 4.0x implied FY27 revenue on two consecutive weekly closes — below this the market prices)
- filing watch: 8-K until 2027-03-09 (deck-provisional: next earnings release, guidance revision, or further strategy/personnel change following the Max pivot)
- dated pack TTANI2026-12-03: TTAN: Q3 FY27 print — four gates: cRPO $\geq 20\%$ YoY, GTV $\geq 15\%$, usage revenue $\geq$ flat sequential, NDR still $>110\%$; any two failing = full exit
- dated pack TTANI2026-12-07: TTAN: Q3 FY27 10-Q filing carrying the RPO footnote and processor concentration language (date confirmed=False)
- dated pack TTANI2026-12-18: TTAN: only options expiry that covers the Dec-3 catalyst while keeping the loss branch inside tax year 2026 (date confirmed=True)
- dated pack TTANI2027-03-10: TTAN: Q4/FY27 print — T3 confirmation add requires positive sequential usage revenue here (date confirmed=False)

What remains unverified

1. **Q3 FY27 usage revenue and GTV trajectory.** Time-gated to the print; genuinely unknowable now. *Estimate:* usage revenue **\$64–70M, i.e. -3% to -11% sequential**, from our own beat-adjusted decomposition net of professional-services runoff and the one business day. Probability it prints **flat or better sequentially: 0.30**.
2. **Current diluted share count including Class C.** Drainable today from the 10-Q cover page and DEI tags; we simply have not pulled it. Every valuation number here is hostage to it. *Working assumption:* 95.4M, with a plausible 95–98M band.
3. **Exact cRPO percentage** (the “50%” and “53%” are rounded). Drainable today from the 10-Q RPO footnote. This determines whether the +23.1% figure is really +20.8% or +25.3% — and therefore whether the bear’s backlog argument is actually dead.
4. **Payment-processor identity, take rate, contract term and renewal date. Not drainable at any price** — private bilateral contract. *Estimate:* the concentration tag implies $>10\%$ of revenue; we assume **60–80% of the \$72.1M usage line** routes through one processor at a **0.4–0.9% take on GTV**, benchmarked to integrated-payments comparables. Probability of a material repricing in the next twelve months: **0.15**.
5. **Average daily volume and our real liquidity footprint.** No price bars on either bench. *Estimate:* **2.5–4.0M shares/day (\$140–225M notional)**, from the base rate for a \$5.4B recent-IPO software name in the month after a $\sim 30\%$ gap. Confidence 0.70.
6. **Austin and Bernalillo County trade-permit series, 2025 vs 2026.** Drainable today via the public Socrata endpoints by work class; no company name needed.
7. **Street consensus for Q3 FY27.** Drainable today; neither bench pulled it. Until we have it, “the market overreacted” is an opinion.
8. **Q1 FY27 subscription-versus-usage split.** Drainable from the Q1 10-Q already on file; it would let us build a three-quarter usage trend instead of a two-point line.

9. **Rule 144 share counts and 10b5-1 plan adoption dates.** Drainable from the broker filer's EDGAR history. Determines whether August's supply is routine or discretionary.
10. **Net cash and debt.** Drainable from the 10-Q balance sheet. Our 4.70x is price-to-sales, not enterprise-value-to-sales; with a meaningful net cash balance the real multiple is lower and the "not cheap" objection weakens.
11. **Confirmation of the 2026-12-03 print date.** Not drainable until the company issues a release. *Estimate: 2026-12-03 ± 7 sessions*, confidence **0.65**, from last year's early-December cadence.
12. **Actual retention value inside the ">110%" band.** Not drainable — the company does not disclose it. *Estimate: 110–116%*, with probability it is **below 114%: 0.55**, inferred from the usage decline and the three-quarter duration slide.

PRIZE TABLE + VOL READ

Items 1, 4, 11 and 12 are genuinely time-gated or private, so the table is owed. **Anchor:** spot \$56.285 = 4.70x implied FY27 revenue. Adjustments are made off that anchor, not off a blank sheet.

#	Scenario	What has to be true	p	Multiple	Price
1	Usage reaccelerates, all four gates pass	December usage flat-to-up sequentially, GTV ≥15%, cRPO holds ~23%, habitual beat returns	0.25	6.5x	\$77.78
2	Muddle through	Guide met, usage down mid-single-digits, retention band holds, multiple unchanged-to-slightly-better	0.35	5.2x	\$62.22
3	Usage break confirmed	Second consecutive negative usage quarter, exit rate settles 13–14%, AI pivot still pre-revenue	0.30	3.5x	\$41.88
4	Severe	Processor repricing or a GTV break, retention drops out of the band, pivot visibly failing	0.10	2.8x	\$33.50

EV = \$57.14, or +1.5% versus spot. That is the number that decides this deck. It is not an edge; it is a rounding error. **This is why there is no immediate entry and why the size is 0.60% and not 1.5%.** 40% of probability mass sits below spot. The whole case for participating at all is that scenario 1 pays 38% and scenario 3 costs 26% — an acceptable shape, at a small size, entered lower.

Dispersion: $\sigma_{\text{table}} = \$15.53 = 27.6\% \text{ of spot}$, over $T = 99$ days to the December monthly expiry ($\sqrt{T} = 0.521$). Our table therefore implies an equivalent annualized volatility of $27.6\% / 0.521 = \sim 53\%$.

Bimodality: the table is fork-dated on 2026-12-03. Almost all of the distance between scenario 1 (\$77.78) and scenario 3 (\$41.88) resolves in a single session. A lognormal volatility number understates the fork and overstates the chance of finishing near the middle. Any premium sold across that date is short a gap, not short a drift.

Expression math — and what we cannot complete. Pricing candidate calls straight off the table ($C_{\text{ours}} = \sum p_i \times \max(S_i - K, 0)$):

Strike (Dec-2026)	C_ours from the table	% of spot	Chain midpoint	Verdict
\$60	\$5.22	9.3%	NOT RETRIEVED	cannot decide
\$65	\$3.20	5.7%	NOT RETRIEVED	cannot decide
\$70	\$1.95	3.5%	NOT RETRIEVED	cannot decide

No options chain was retrieved on either bench, and no price-bar series exists to compute realized volatility. We therefore cannot compute the variance risk premium (IV minus RV) and we will not fake one. The decision rule is written out so it can be executed the moment the chain arrives:

- If **chain-level midpoint IV** on the Dec-18-2026 \$60 and \$65 strikes is **below ~53%**, our table's dispersion exceeds the chain's and the gated T2 add should be expressed in **calls** rather than stock.
- If it is **above ~53%** — which we consider more likely for a name that just gapped roughly 30%, where 60-day realized is plausibly 70–90% — the chain is pricing more dispersion than our own table, calls are the wrong instrument, and the trade routes to the **premium-selling gates** (cash-secured puts struck at or below the \$52.00–54.50 T1 band, which pays us to place the bid we already want to place).
- **Use per-strike chain midpoints and real bid/ask quotes, never the underlying's aggregate IV feed.** The TSSI lesson stands: the aggregate feed read 71 while the actual chain read 101. They diverge, and the divergence is exactly the size of the decision.

Tax overlay. The loss branch of this table (scenarios 3 and 4, 40% of the mass) resolves at the December 3 print. In a harvest year that loss must land in **tax year 2026**. Stock entered now satisfies this — we can realize before December 31. If the trade is later expressed in options, **only the 2026-12-18 expiry qualifies**: it covers the December 3 catalyst *and* the loss branch settles inside 2026. **Any January or February 2027 expiry defers the loss into the following year and**

is excluded during a harvest year, notwithstanding that the later expiry buys more time for scenario 1 — the edge here (+1.5% EV) comes nowhere close to numerically dominating the deferral cost, so the exception does not apply.

Instrument recommendation — the output of the arithmetic above: COMMON STOCK, staged, at the T1 band. Not calls, because we cannot show they are cheap without the chain and the fork argues the chain is expensive. Not covered calls, because we have no shares yet. Not cash-secured puts *yet*, because pricing them without real quotes would be guessing. Not “none,” because the scenario shape is acceptable at small size. **The chain pull is the single highest-value drain in this deck after the permit series, and it flips the instrument the day it lands.**

```

{"price_gates": [{"level": 54.5, "direction": "below", "basis": "T1 starter band top: 4.55x implied FY27 revenue, at the 52-week low of 54.17 – bid into weakness, limit orders only"}, {"level": 54.17, "direction": "below", "basis": "52-week low from the pack tape; a decisive break reclassifies this from un-absorbed knife to new leg down"}, {"level": 47.86, "direction": "below", "basis": "hard stop at 4.0x implied FY27 revenue on two consecutive weekly closes – below this the market prices a sub-12% grower and the thesis is wrong"}, {"level": 71.79, "direction": "above", "basis": "trim half at 6.0x if usage revenue is still negative sequentially – an unearned multiple re-rate"}], "event_gates": [{"forms": ["8-K"], "why": "next earnings release, guidance revision, or further strategy/personnel change following the Max pivot and the CRO change"}, {"forms": ["10-Q"], "why": "RPO footnote with the exact next-12-month percentage, subscription-vs-usage split, and any text change to the third-party processor concentration language"}, {"forms": ["4", "144"], "why": "August 2026 ran four 144s in eleven days plus three Form 4s – track whether insider supply into the tape persists"}, {"forms": ["S-1", "S-3", "424B4"], "why": "a follow-on would add supply on top of the existing 144 cadence at a 52-week low"}], "catalyst_dates": [{"date": "2026-12-03", "what": "Q3 FY27 print – four gates: cRPO >=20% YoY, GTV >=15%, usage revenue >= flat sequential, NDR still >110%; any two failing = full exit", "confirmed": false}, {"date": "2026-12-07", "what": "Q3 FY27 10-Q filing carrying the RPO footnote and processor concentration language", "confirmed": false}, {"date": "2026-12-18", "what": "only options expiry that covers the Dec-3 catalyst while keeping the loss branch inside tax year 2026", "confirmed": true}, {"date": "2027-03-10", "what": "Q4/FY27 print – T3 confirmation add requires positive sequential usage revenue here", "confirmed": false}], "immediate_entry": null, "unverified": [{"item": "Q3 FY27 usage revenue and GTV trajectory – the core open question", "instrument": "2026-12-03 print + 10-Q segment detail", "drainable_now": false, "estimate": {"range": "$64-70M usage revenue, -3% to -11% sequential", "basis": "beat-adjusted guide decomposition net of negative-margin professional-services runoff and one fewer business day (~$1.1M)", "p": 0.3}}, {"item": "Current diluted share count including the Class C line – denominator for every valuation number here", "instrument": "XBRL pull, DEI/cover-page tags on 10-Q 0001638826-26-000094", "drainable_now": true, "estimate": null}, {"item": "Exact cRPO percentage (50% and 53% are rounded; true YoY range is +20.8% to +25.3%)", "instrument": "XBRL pull, RPO footnote of 10-Q 0001638826-26-000094 and IR supplemental", "drainable_now": true, "estimate": null}, {"item": "Third-party payment processor identity, take rate, contract term and renewal date", "instrument": "none – private bilateral contract, no public observable; PRIVATE BILATERAL COUNTERPARTY-TERMS class", "drainable_now": false, "estimate": {"range": "60-80% of the $72.1M usage line via one processor at a 0.4-0.9% take on GTV", "basis": ">10% concentration implied by the disclosure tag; integrated-payments comparables", "p": 0.15}}, {"item": "Average daily volume and our real liquidity footprint versus a 0.60% book position", "instrument": "IBKR historical bars", "drainable_now": false, "estimate": {"range": "2.5-4.0M shares/day, $140-225M notional; our 26.6k shares = 0.7-1.1% of ADV", "basis": "base rate for a $5.4B recent-IPO software name one month after a ~30% gap", "p": 0.7}}, {"item": "Austin and Bernalillo County monthly HVAC/plumbing/electrical permit counts, 2025 vs 2026", "instrument": "City of Austin Socrata dated permit series by work class; Bernalillo County monthly trade permits – no issuer name required", "drainable_now": true, "estimate": null}, {"item": "Street consensus for Q3 FY27 revenue", "instrument": "web_search / consensus connector", "drainable_now": true, "estimate": null}, {"item": "Q1 FY27 subscription-versus-usage revenue split", "instrument": "XBRL pull on 10-Q 0001638826-26-000047 (Q1 FY27, already filed)", "drainable_now": true, "estimate": null}, {"item": "Rule 144 share counts and 10b5-1 plan adoption dates behind the August cluster", "instrument": "EDGAR full-text pull on filer 0001958244, forms 144 dated 2026-08-04/11/12/14", "drainable_now": true, "estimate": null}, {"item": "Net cash and debt – our 4.70x is price-to-sales, not enterprise-value-to-sales", "instrument": "XBRL pull, balance sheet of 10-Q 0001638826-26-000094", "drainable_now": true,

```

```
"estimate": null}, {"item": "Confirmation of the 2026-12-03 Q3 print date",
"instrument": "company press release / IR calendar; 8-K event gate",
"drainable_now": false, "estimate": {"range": "2026-12-03 +/- 7 sessions",
"basis": "prior-year early-December Q3 cadence; current date is yfinance-
derived only", "p": 0.65}}, {"item": "Actual net dollar retention value
inside the '>110%' disclosure band", "instrument": "none – issuer discloses a
band, not a number; BANDED-METRIC MASKING class", "drainable_now": false,
"estimate": {"range": "110-116%", "basis": "implied usage-line decline plus
the 53%/52%/50% three-quarter duration slide", "p": 0.55}}, {"item": "Options
chain: Dec-18-2026 $60 and $65 midpoint IV and real bid/ask quotes – flips
the instrument recommendation", "instrument": "IBKR option chain, per-strike
midpoints only, never the underlying aggregate IV feed", "drainable_now":
true, "estimate": null}, {"item": "Realized volatility series for the IV-
minus-RV comparison", "instrument": "IBKR daily bars, 30/60-day realized",
"drainable_now": false, "estimate": {"range": "60-day RV 70-90% annualized",
"basis": "base rate following a single-session ~30% gap in a $5B software
name", "p": 0.6}}}]}
```

ADJUDICATION — 2026-09-16 15:38Z

REDUCED DAMAGE-ARRIVING: FY guide raised, cRPO +23.1%, and red's strongest kill was a guide-vs-actual base effect; but growth exits H2 at ~14–15% vs 25% four quarters ago and the equity story now leans on an AI-product pivot announced with a CRO change. Software cohort → no adds before 9/21 (regime flow gate). Entry staged, not chased, near the low.

REOPEN GATES: 9/22: stage 0.4x starter (~0.4% NAV) limit at/below the 9/21 close · Q3 print ~12/03: cRPO ≥20% YoY, GTV ≥15%, usage revenue ≥ flat sequential, NDR >110% — any two failing → exit, not add

PROCESS: austin/albuquerque permits unrun — highest-value test of the GTV leg (ticket)